TruthGuard: AI-Powered Fake News Detection & Verification Tool

Name: Aditya Mukeshkumar Gupta   Batch 10

Project Statement: Addressing the Global Crisis of Misinformation with TruthGuard

The proliferation of misinformation, often referred to as "fake news," across global social media platforms, messaging apps, and web channels represents one of the most critical threats to modern society. This rapid, often viral, dissemination erodes public trust in traditional media and institutions, poses severe risks to public safety (especially in health crises like pandemics), influences democratic processes through targeted disinformation campaigns, and contributes to societal polarization.

The Existing Challenge and Critical Gap:

Current efforts to combat this crisis are fundamentally insufficient. The sheer volume and velocity of new content published every minute far outstrips the capacity of human fact-checkers. Manual fact-checking is an inherently time-consuming process, requiring researchers to identify a claim, locate primary sources, verify evidence, and write a detailed debunking article. By the time a piece of misinformation is manually verified as false, it has often already reached millions of users and caused irreparable harm. Consequently, the existing manual approach is neither a real-time solution nor is it scalable to meet the demands of the modern digital information ecosystem.

Introducing TruthGuard: A Unified, Automated, and Explainable Solution

TruthGuard is designed to fill this critical operational gap by creating a unified, end-to-end automated platform for real-time news authenticity classification. The core objective is to move beyond simple binary (True/False) classification and provide a granular assessment of information integrity.

The system's primary function is to classify any input news article, social media post, or textual claim into one of three distinct and actionable categories:

1. **Real (Authentic):** Content verified to be factually accurate and sourced credibly.
2. **Suspicious (Requires Caution):** Content containing sensational language, misleading framing, unverified claims, or partially false information (misinformation).
3. **Fake (Disinformation):** Content determined to be deliberately fabricated, contextually manipulated, or entirely false.

Technological Foundation and Methodology:

TruthGuard leverages state-of-the-art advancements in computational linguistics and artificial intelligence to achieve its objectives:

- **Natural Language Processing (NLP) and Deep Learning:** The system employs sophisticated NLP techniques to analyze both the linguistic features (e.g., tone, emotional charge, use of hyperbole) and structural elements of the text.
- **Transformer Models (BERT/RoBERTa):** At the heart of the classification engine are pre-trained Transformer models, such as BERT (Bidirectional Encoder Representations from Transformers) and RoBERTa (Optimized Robustly Optimized BERT Pretraining Approach). These models are fine-tuned on massive datasets of verified and fake news to understand complex semantic relationships, identify subtle manipulative language patterns, and detect stylistic cues associated with false narratives.
- **Hybrid Verification Pipeline:** To provide robust verification, the system does not rely solely on internal textual analysis. It simultaneously verifies extracted claims and entities against trusted external fact-checking databases, journalistic archives, and reputable academic APIs (where available). This hybrid approach significantly increases the system's accuracy and robustness against novel disinformation tactics.

The integration of these technologies results in a powerful tool that offers both explainability (showing *why* a piece of news was flagged as Suspicious or Fake by highlighting suspicious phrases or failed verification points) and scalability, ensuring that the fight against misinformation can finally keep pace with its real-time spread.

Outcomes

The successful implementation of this project delivers the following key outcomes:

- **Automated Credibility Assessment:** Instant detection of fake news with a probability score (0–1) and classification (Reliable/Suspicious/Fake).
- **Real-Time Verification:** Automatic cross-referencing of claims using trusted external APIs like Google Fact Check and PolitiFact.
- **Explainable AI:** Users receive detailed insights, including highlighted suspicious phrases and source credibility scoring, rather than just a simple label.
- **High Performance:** An ultra-fast analysis response (aiming for <100ms) achieved through a hybrid approach of rule-based detection and AI models.
- **Interactive Dashboards:** A user-facing interface for analysis and an administrative dashboard for monitoring system health and user activity.

Outcomes and Deliverables

The successful implementation of this sophisticated project is designed to deliver a suite of impactful and measurable outcomes, fundamentally enhancing the capability to combat the spread of misinformation and disinformation. The key deliverables are structured around speed, accuracy, transparency, and user experience.

Outcome Category	Key Deliverable	Detailed Description
Core Detection	**Automated Credibility Assessment**	The system will provide an instant, data-driven analysis of news articles and claims. This includes a quantitative probability score ranging from 0 to 1, indicating the likelihood of the content being factual. Based on this score, the content will be classified into three distinct, user-friendly categories: Reliable, Suspicious, or Fake. This automated classification minimizes human intervention and speeds up the initial triage of information.

Verification & Trust	**Real-Time Verification Engine**	To ensure maximum accuracy, the system incorporates an intelligent real-time verification module. It automatically cross-references key claims, entities, and data points within the analyzed text against established and trusted external fact-checking databases and APIs. This includes leveraging authoritative sources such as the Google Fact Check API and well-regarded organizations like PolitiFact, providing external validation and bolstering the system's own assessment.
Transparency & Adoption	**Explainable AI (XAI) Module**	Moving beyond a simple 'black-box' label, this outcome focuses on building user trust through transparency. Users will receive detailed, actionable insights into *why* a particular credibility assessment was reached. This includes:Highlighting specific suspicious phrases or linguistic patterns that triggered the system's flags.Providing a comprehensive source credibility scoring, analyzing the historical accuracy, publication standards, and editorial policies of the originating source.
System Efficiency	**High-Performance Analysis Architecture**	Recognizing the need for rapid response in a fast-moving information environment, the system is engineered for ultra-low latency. The core objective is to achieve an analysis and response time of less than 100 milliseconds (<100ms) for typical-length articles. This high performance is achieved through a hybrid detection approach that efficiently combines the speed of rule-based pattern matching and natural language processing (NLP) with the deep inference capabilities of advanced machine learning and AI models.
User Interfaces	**Interactive & Intuitive Dashboards**	The project includes the development of two specialized, role-based interfaces:User-Facing Analysis Interface: A clean, intuitive dashboard where end-users can submit links or text for analysis and receive instant, explainable results, fostering direct engagement with the verification process. Administrative Monitoring Dashboard: A secure backend interface for system administrators and data scientists, providing a holistic view of system health, real-time performance metrics, model drift warnings, and aggregated user activity/usage statistics, ensuring continuous optimization and maintenance.

Comprehensive System Architecture: Detailed Module Breakdown

The FastNewsDetector system is architecturally segmented into seven interconnected and specialized modules, each responsible for a critical phase of the news analysis and verification workflow. This modular approach ensures scalability, maintainability, and clear separation of concerns.

Module 1: User Authentication and Security

This foundational module manages all aspects of user identification and access control, ensuring a secure environment for analysis.

- **User Registration and Profile Management:** Facilitates secure sign-up of new users, collecting essential profile information, and allowing existing users to manage their credentials.
- **User Login and Session Management:** Implements robust login mechanisms and manages active user sessions efficiently using Flask-Login. This ensures statefulness and personalization throughout the user experience.
- **Password Security:** Adopts industry best practices by utilizing PBKDF2 (Password-Based Key Derivation Function 2) for one-way hashing and salting of all stored passwords, making them resistant to brute-force and rainbow table attacks.
- **API and Inter-Service Security:** Employs JSON Web Tokens (JWT) for securing APIs and internal communications between microservices, ensuring that only authenticated and authorized requests are processed.

Module 2: Input Processing and Acquisition

The entry point of the analysis pipeline, this module is responsible for accepting, validating, and standardizing the source material.

- **Flexible Input Acceptance:** The system accepts two primary forms of input:
    a. **Raw Text:** Direct copy-pasted articles or snippets.
    b. **URLs:** Links to online articles, blog posts, or social media content.
- **Web Scraping and Parsing:** When a URL is provided, the module uses sophisticated web scraping techniques, leveraging the BeautifulSoup library, to intelligently extract the

main textual content while discarding boilerplate navigation, advertisements, and irrelevant HTML elements.

- **Input Validation:** Prior to initiating resource-intensive analysis, the module performs crucial checks:
    a. **Length Validation:** Ensures the text is within a reasonable limit (e.g., minimum word count to allow for meaningful analysis, and a maximum to prevent Denial-of-Service attacks).
    b. **Format and Encoding:** Confirms proper text encoding (e.g., UTF-8) and basic structural integrity.

Module 3: NLP & Feature Engineering Pipeline

The core intelligence layer where raw text is transformed into quantifiable, machine-readable features essential for classification.

- **Advanced Text Preprocessing:** The text undergoes a rigorous cleaning process:
    a. **Tokenization:** Breaking the text into individual words or sub-word units.
    b. **Stopword Removal:** Eliminating common, low-information words (e.g., "the," "a," "is") using curated lists from NLTK/spaCy.
    c. **Lemmatization/Stemming:** Reducing words to their base or root form (Lemmatization via spaCy is preferred for retaining context) to minimize feature space and improve model accuracy.
- **Feature Extraction:** Generating contextual and structural indicators:
    a. **Sentiment Intensity Score:** Utilizing specialized lexical analysis tools (e.g., VADER or similar) to quantify the overall emotional tone (positive, negative, neutral) and intensity of the article.
    b. **Sensationalism and Hyperbole Score:** Calculating metrics based on the frequency of all-caps words, excessive punctuation (e.g., multiple exclamation marks), hyperbolic adjectives, and emotionally charged vocabulary, indicative of clickbait or sensational reporting.

Module 4: Dual-Layer Classification Engine

This module contains the primary logic for determining the veracity and quality of the news content, employing a synergistic two-pronged approach.

- **Layer 1: FastNewsDetector (Rule-Based Expert System):**
  a. **Purpose:** Provides an immediate, low-latency preliminary assessment. Ideal for flagging obvious indicators of poor quality or misleading content.
  b. **Mechanism:** A deterministic system relying on custom regex (regular expressions) patterns to detect common red flags (e.g., misleading domain names, highly inflammatory phrases, presence of known manipulative terms) and **indicator counting** (e.g., number of citations, ratio of facts to opinion).
- **Layer 2: GeminiAssistant (AI-Powered Deep Analysis):**
  a. **Purpose:** Executes a deep, contextual, and semantic analysis for complex or nuanced cases that escape simple rule-based detection.
  b. **Mechanism:** Integration with the Google Gemini API (or equivalent large language model) to perform:
     i. **Semantic Consistency Checks:** Assessing the logical flow and self-contradictions within the text.
     ii. **Contextual Verification:** Evaluating claims against general world knowledge and recent events.
     iii. **Detailed Explanations:** Generating human-readable summaries justifying the classification decision.

Module 5: Claim Verification and External Fact-Checking

This critical module acts as the external validation bridge, cross-referencing specific, verifiable statements against trusted external databases.

- **Claim Identification:** Utilizing Named Entity Recognition (NER) and dependency parsing (from **spaCy**) to automatically isolate explicit, verifiable claims (e.g., "The president said X," "The company's stock rose by Y percent").
- **API Integration:** Seamlessly integrates with established, third-party fact-checking APIs (e.g., Poynter's International Fact-Checking Network database or dedicated news verification services).

- **Cross-Referencing and Score Aggregation:** Submitting identified claims to external databases and aggregating the verification results, assigning a confidence score or verification status (e.g., True, False, Unproven, Mixture) to the specific claim.

Module 6: Explainability, Reporting & Visualization

This user-facing module translates complex analytical data into actionable, easily understandable insights.

- **Suspicious Text Highlighting:** Provides visual feedback by dynamically highlighting specific sentences or phrases in the input text that triggered high scores in the sensationalism or misclassification detectors (e.g., color-coding based on severity).
- **Card-Based User Interface (UI):** Presents the aggregated findings in a clear, modern, and intuitive card-based layout. Each card represents a key finding:
    a. Overall Credibility Score.
    b. Sentiment Analysis Breakdown.
    c. Fact-Check Verification Status (from Module 5).
    d. Summary of GeminiAssistant's semantic analysis.
- **Transparency and Trust:** Focuses on *why* a decision was made, not just *what* the decision was, enhancing user trust in the system's output.

Module 7: Administrative and Operational Dashboard

A dedicated module for system administrators to monitor performance, manage resources, and ensure the continuous health of the service.

- **System Usage Statistics:** Provides high-level metrics, including:
    a. Total number of analyses performed globally and over time.
    b. Breakdown of input types (Text vs. URL).
    c. Peak usage hours and geographical distribution.
- **User Activity Monitoring:** Tracks active users, new registrations, and specific module usage for identifying popular features and potential abuse.
- **System Health and Resource Monitoring:** Offers real-time visibility into the operational status of critical components:
    a. **Database Performance:** Query times and connection status.

b. **Detector Speed:** Latency metrics for the rule-based and AI classification layers.

c. **API Status:** Up/Down status and rate-limit consumption for all third-party and internal APIs (Gemini, Fact-Check APIs).

Architecture Diagram

The system follows a Layered Architecture designed for scalability and separation of concerns.

- **Presentation Layer:** Flask HTML Templates (Jinja2) & JSON APIs for the frontend interface.
- **Business Logic Layer:** Contains the core intelligence, including the `FastNewsDetector` class for speed and `GeminiAssistant` class for AI analysis.
- **Data Access Layer:** Uses SQLAlchemy ORM to manage database interactions and queries.
- **Storage Layer:** SQLite database for persistent data storage and a file system for logs.

Data Flow:

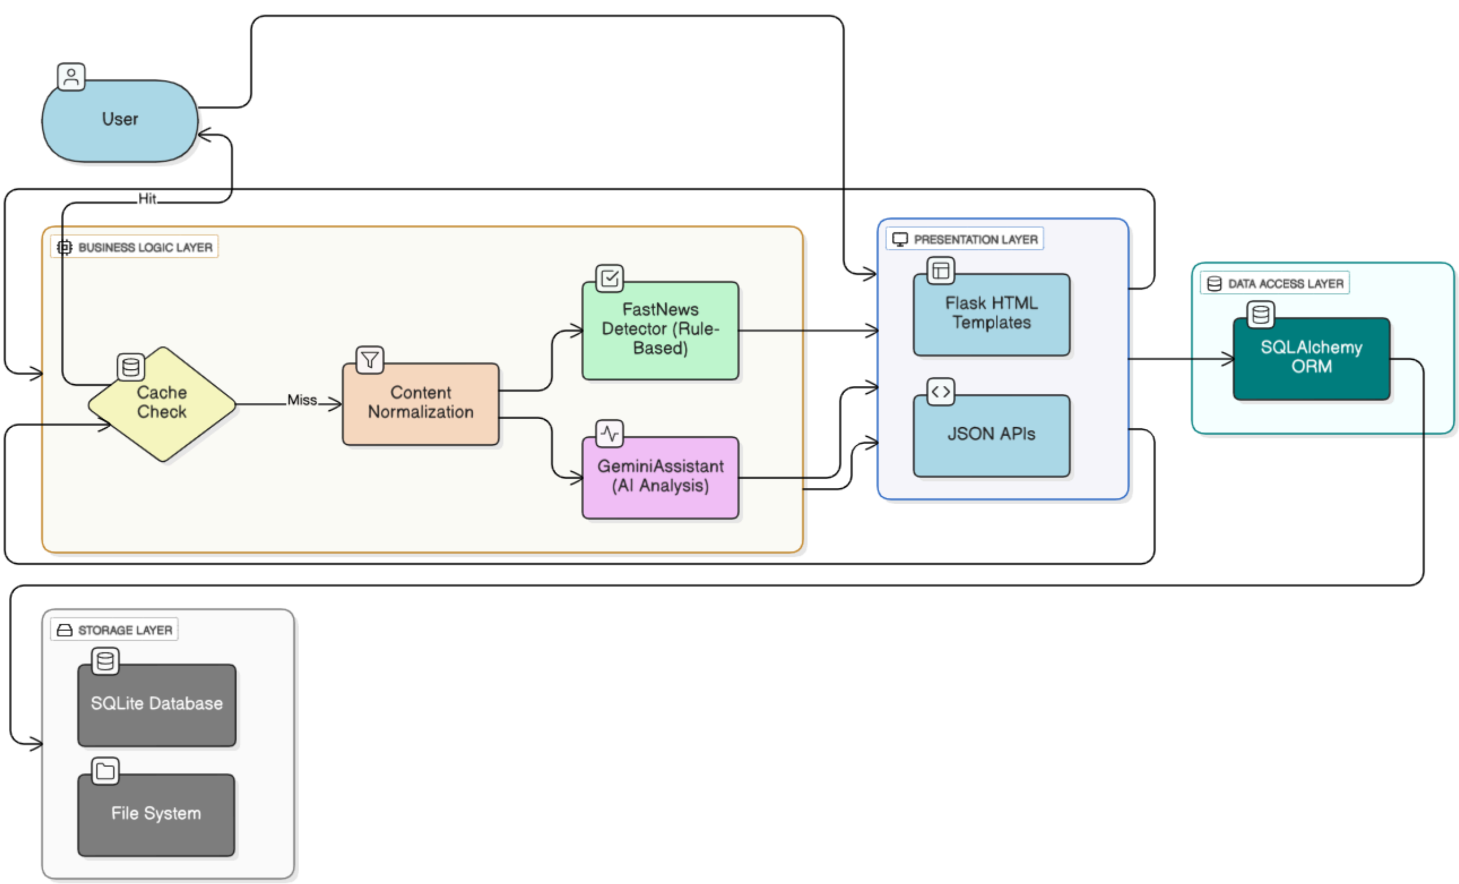

Database Schema: Detailed Structure and Purpose

The application leverages a robust, relational database architecture, primarily managed using SQLite for development and potentially PostgreSQL/MySQL in production, interfaced through the SQLAlchemy ORM. This design ensures data integrity, scalability, and efficient querying. The database is logically segmented into several key tables, each serving a distinct functional purpose within the fake news detection and analysis ecosystem:

A. Users Table (User)

This is the core table for managing all registered individuals within the system, crucial for authentication, authorization, and tracking user-generated content.

Field Name	Data Type	Constraints	Description
id	Integer	Primary Key	Unique identifier for each user.
email	String (255)	Unique, Not Null	The user's login ID and primary contact point.
password_hash	String (255)	Not Null	A securely hashed representation of the user's password using a robust algorithm like bcrypt.
name	String (100)	Not Null	The full or display name of the user.
role	Enum	Not Null, Default: '…	Defines the user's access level and privileges (e.g., **Admin**, **Moderator**, **User**).
created_at	DateTime	Not Null	Timestamp recording the exact moment the user account was created.
last_login	DateTime	Nullable	Timestamp of the user's most recent successful login.

B. Articles/Analysis Table (`Analysis`)

This table is central to the application, storing the core content (articles, social media posts, etc.) that have been submitted for verification, along with the results of the automated and manual analysis.

Field Name	Data Type	Constraints	Description
id	Integer	Primary Key	Unique identifier for the analysis record.
user_id	Integer	Foreign Key (Re…	The ID of the user who submitted the article for analysis. Null if submitted by an automated process.
content	Text	Not Null	The full body or a substantial excerpt of the text being analyzed.
classification	Enum	Not Null	The final veracity judgment: RELIABLE, SUSPICIOUS (requires further review), or FAKE.
confidence_score	Float	Not Null, Range…	A quantitative measure of the AI model's certainty in its classification.
sensationalism_score	Float	Not Null, Range…	A metric indicating the level of emotionally charged or exaggerated language used in the content.
key_findings	JSON	Nullable	A structured data field containing the AI's most important extracted elements, such as named entities, links to sources, or quotes flagged for verification.
submitted_at	DateTime	Not Null	Timestamp of when the analysis request was received.
processed_duration	Integer	Nullable	Time in seconds taken to complete the analysis.

C. Chat History Table (`GeminiChat`)

This table maintains a record of all interactions with the "TruthBot" AI assistant, ensuring context retention across sessions and providing a valuable dataset for usage analysis.

Field Name	Data Type	Constraints	Description
id	Integer	Primary Key	Unique identifier for the chat session record.
user_id	Integer	Foreign Key (R…	The ID of the user involved in the conversation.
session_id	String (50)	Not Null	A unique identifier linking all messages within a single conversational thread.
message_content	Text	Not Null	The actual text of the message (either user query or AI response).
is_user_message	Boolean	Not Null	Flag indicating if the message originated from the user (`True`) or the AI (`False`).
timestamp	DateTime	Not Null	The time the message was sent/received.
metadata	JSON	Nullable	Stores contextual information, such as the specific model used (e.g., `Gemini-Pro`), token usage, or related analysis IDs.

D. Verification Results Table (`VerificationResult`)

This table stores the factual data retrieved from external, authoritative fact-checking services, allowing the application to cross-reference and validate its own internal analysis.

Field Name	Data Type	Constraints	Description
id	Integer	Primary Key	Unique identifier for the verification result.
analysis_id	Integer	Foreign Key (Ref…	Links the external result back to the specific content being analyzed.
source_api	String (100)	Not Null	The name of the external fact-checking API used (e.g., `FactCheck.org`, `Snopes`).
query_text	Text	Not Null	The specific phrase or claim sent to the external API.
external_rating	String (50)	Not Null	The veracity rating provided by the external source (e.g., 'False', 'Mostly True', 'Unproven').
external_url	String (500)	Nullable	A direct link to the external fact-check article for user review.
checked_at	DateTime	Not Null	Timestamp of when the external verification check was performed.

Conclusion

AI-powered Fake News Detection is an essential tool in the modern digital age. By combining traditional NLP techniques with advanced Deep Learning (Transformers) and real-time fact verification, TruthGuard ensures content integrity and helps build a safer, more informed digital society. The project successfully demonstrates how a hybrid architecture can provide both the speed required for user retention and the accuracy needed for trust.

Sample Output

The system provides immediate JSON responses for API integration and visual cards for the web interface.

Sample JSON Response (API):


```
{
  "success": true,
  "classification": "FAKE",
  "confidence": 91.3,
  "sensationalism": 8.5,
  "credibility": 2.1,
  "key_findings": [
    "Contradicts verified sources regarding this event.",
    "Uses highly emotional and sensationalist language."
  ],
  "processing_ms": 85.2,
  "is_quick": true
}
```

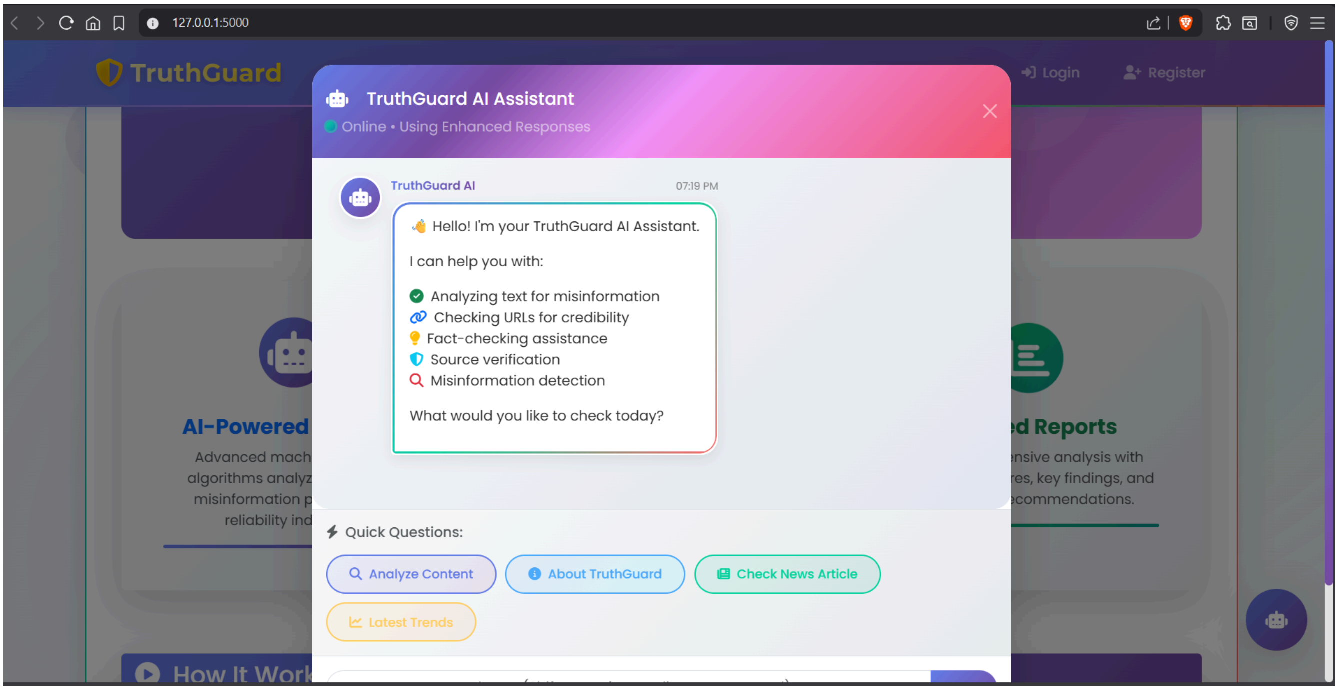
127.0.0.1:5000
TruthGuard
Login
Register
TruthGuard AI Assistant
Online • Using Enhanced Responses
TruthGuard AI
07:19 PM
Hello! I'm your TruthGuard AI Assistant.
I can help you with:
Analyzing text for misinformation
Checking URLs for credibility
Fact-checking assistance
Source verification
Misinformation detection
What would you like to check today?
Quick Questions:
Analyze Content
About TruthGuard
Check News Article
Latest Trends
AI-Powered
Advanced mach
algorithms analyz
misinformation p
reliability ind
d Reports
nsive analysis with
res, key findings, and
ecommendations.
How It Work

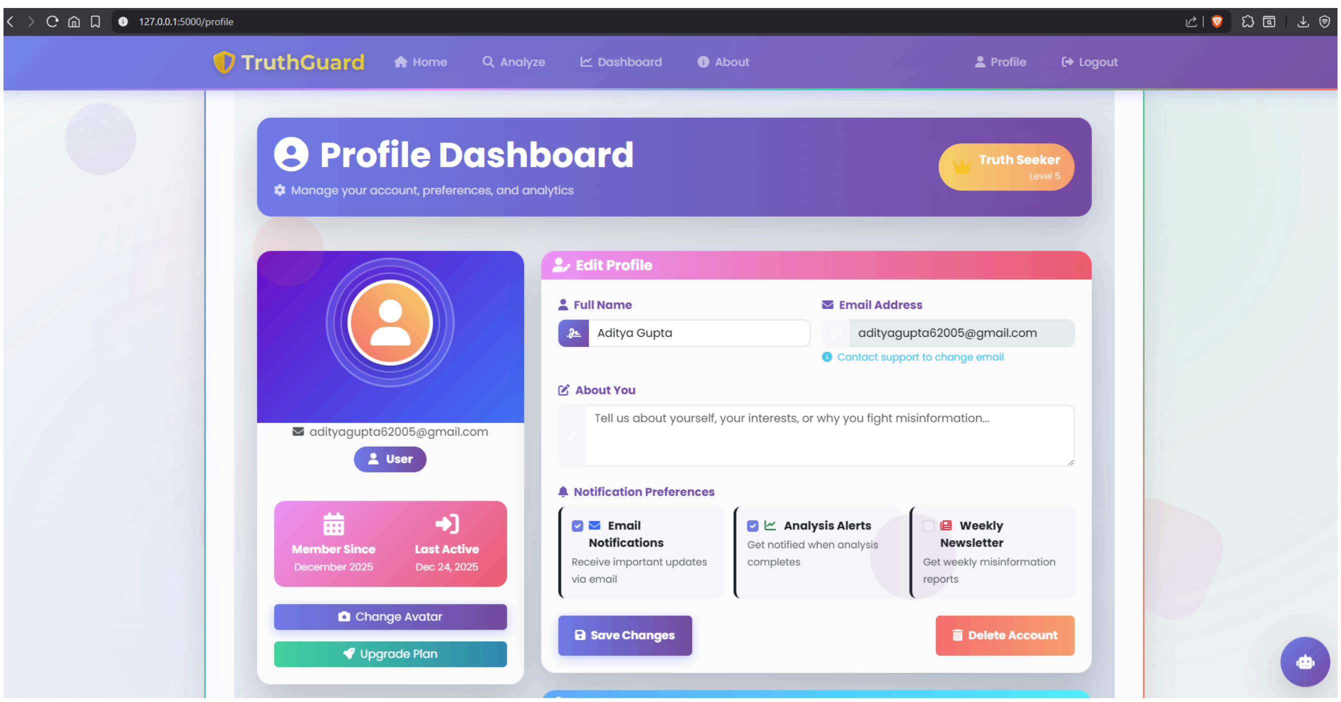
127.0.0.1:5000/profile
TruthGuard
Home
Analyze
Dashboard
About
Profile
Logout
Profile Dashboard
Manage your account, preferences, and analytics
Truth Seeker
Level 5
adityagupta62005@gmail.com
User
Member Since
December 2025
Last Active
Dec 24, 2025
Change Avatar
Upgrade Plan
Edit Profile
Full Name
Aditya Gupta
Email Address
adityagupta62005@gmail.com
Contact support to change email
About You
Tell us about yourself, your interests, or why you fight misinformation...
Notification Preferences
Email Notifications
Receive important updates via email
Analysis Alerts
Get notified when analysis completes
Weekly Newsletter
Get weekly misinformation reports
Save Changes
Delete Account

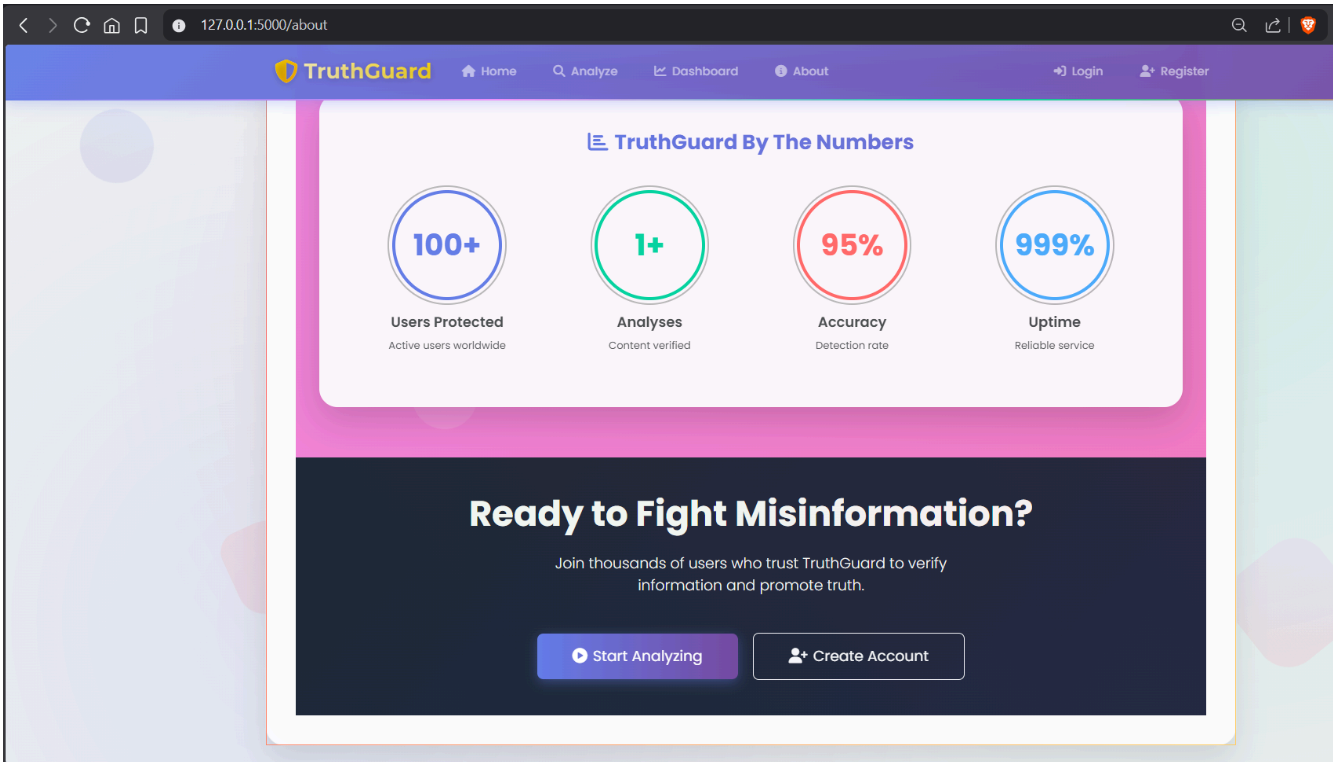
127.0.0.1:5000/about
TruthGuard
Home
Analyze
Dashboard
About
Login
Register
TruthGuard By The Numbers
100+
Users Protected
Active users worldwide
1+
Analyses
Content verified
95%
Accuracy
Detection rate
999%
Uptime
Reliable service
Ready to Fight Misinformation?
Join thousands of users who trust TruthGuard to verify information and promote truth.
Start Analyzing
Create Account

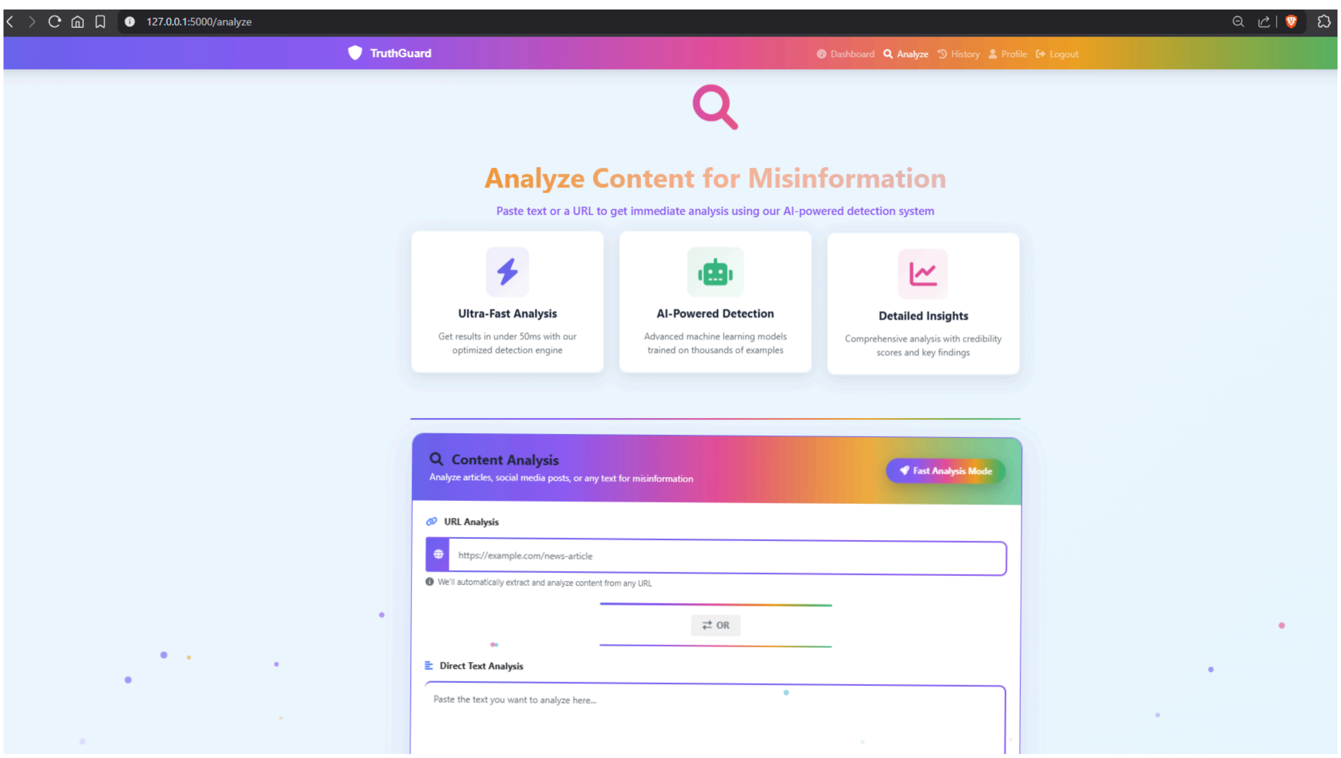
127.0.0.1:5000/analyze
TruthGuard
Dashboard
Analyze
History
Profile
Logout
Analyze Content for Misinformation
Paste text or a URL to get immediate analysis using our AI-powered detection system
Ultra-Fast Analysis
Get results in under 50ms with our optimized detection engine
AI-Powered Detection
Advanced machine learning models trained on thousands of examples
Detailed Insights
Comprehensive analysis with credibility scores and key findings
Content Analysis
Analyze articles, social media posts, or any text for misinformation
Fast Analysis Mode
URL Analysis
https://example.com/news-article
We'll automatically extract and analyze content from any URL
OR
Direct Text Analysis
Paste the text you want to analyze here...

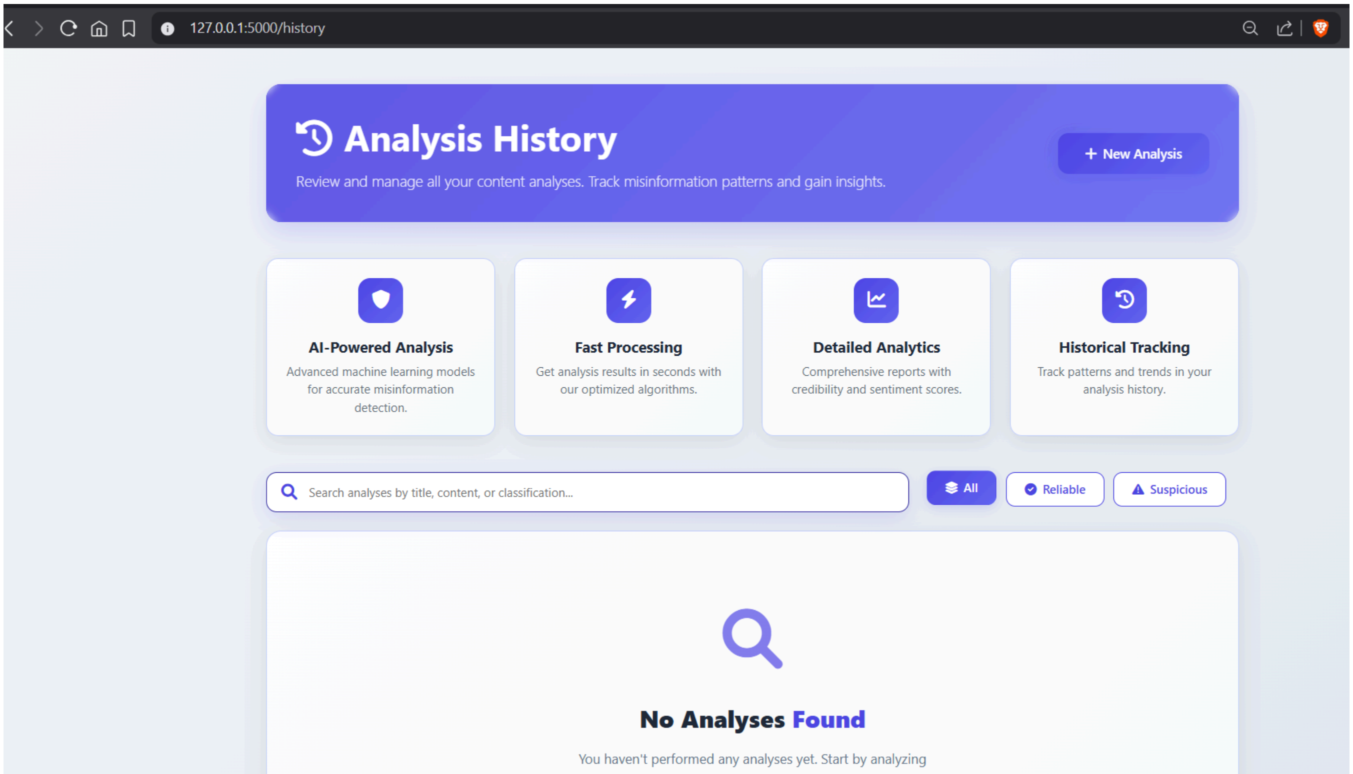
127.0.0.1:5000/history
Analysis History
Review and manage all your content analyses. Track misinformation patterns and gain insights.
+ New Analysis
AI-Powered Analysis
Advanced machine learning models for accurate misinformation detection.
Fast Processing
Get analysis results in seconds with our optimized algorithms.
Detailed Analytics
Comprehensive reports with credibility and sentiment scores.
Historical Tracking
Track patterns and trends in your analysis history.
Search analyses by title, content, or classification...
All
Reliable
Suspicious
No Analyses Found
You haven't performed any analyses yet. Start by analyzing

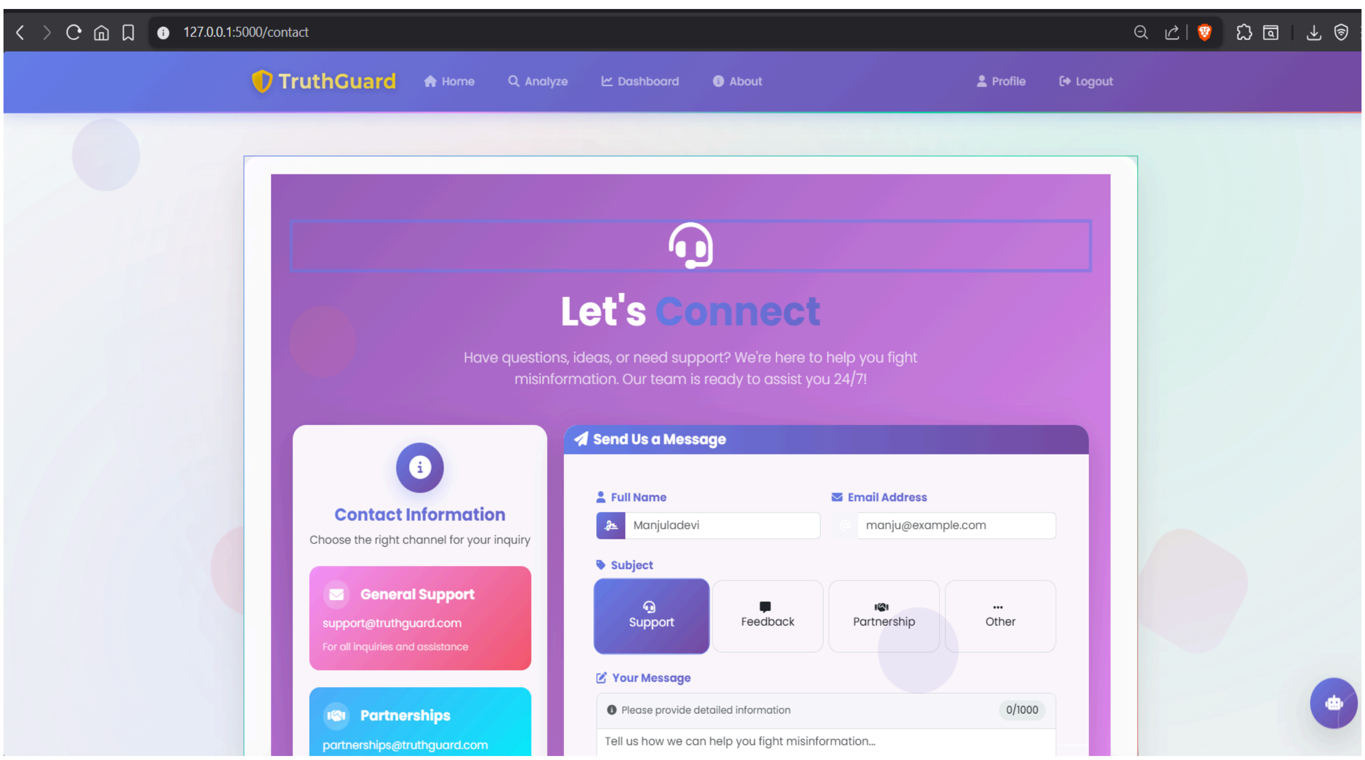
127.0.0.1:5000/contact
TruthGuard
Home
Analyze
Dashboard
About
Profile
Logout
Let's Connect
Have questions, ideas, or need support? We're here to help you fight misinformation. Our team is ready to assist you 24/7!
Contact Information
Choose the right channel for your inquiry
General Support
support@truthguard.com
For all inquiries and assistance
Partnerships
partnerships@truthguard.com
Send Us a Message
Full Name
Manjuladevi
Email Address
manju@example.com
Subject
Support
Feedback
Partnership
Other
Your Message
Please provide detailed information
0/1000
Tell us how we can help you fight misinformation...

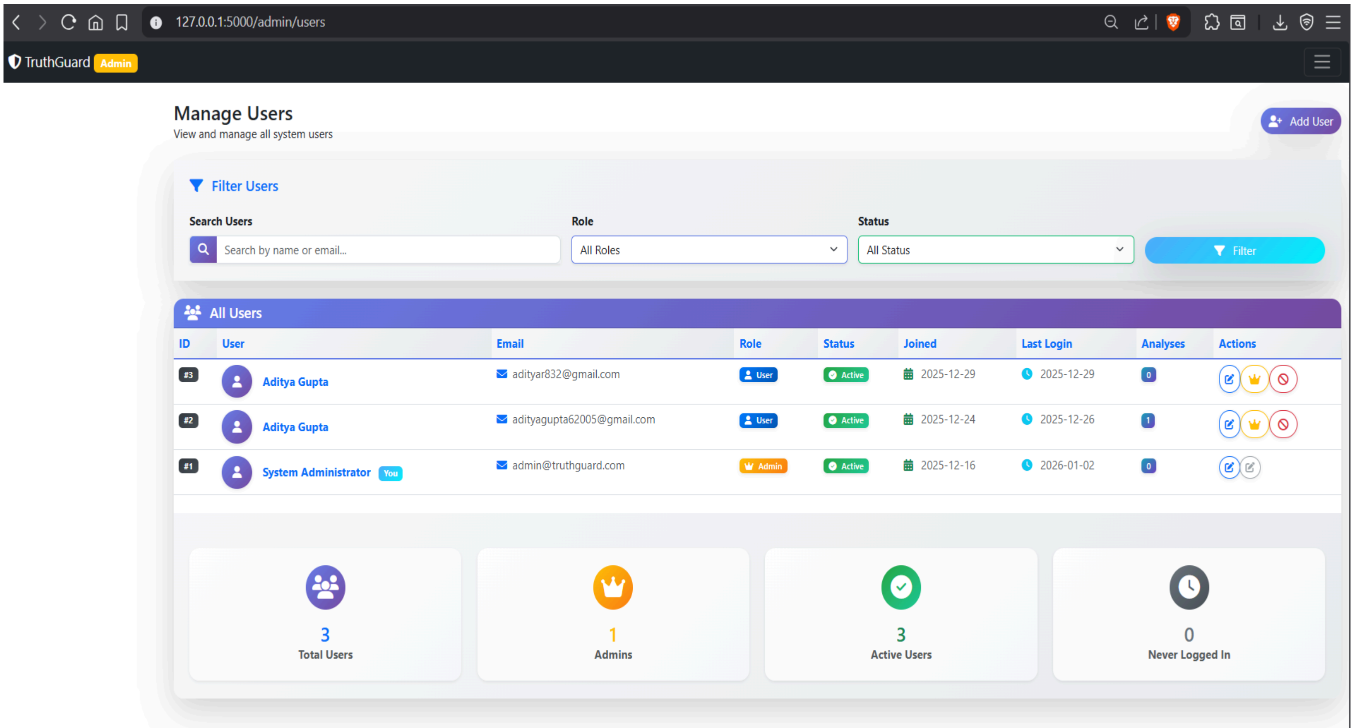

127.0.0.1:5000/admin/users
TruthGuard Admin
Manage Users
View and manage all system users
Add User
Filter Users
Search Users
Search by name or email...
Role
All Roles
Status
All Status
Filter
All Users
ID User Email Role Status Joined Last Login Analyses Actions
#3 Aditya Gupta adityar832@gmail.com User Active 2025-12-29 2025-12-29 0
#2 Aditya Gupta adityagupta62005@gmail.com User Active 2025-12-24 2025-12-26 1
#1 System Administrator You admin@truthguard.com Admin Active 2025-12-16 2026-01-02 0
3 Total Users
1 Admins
3 Active Users
0 Never Logged In

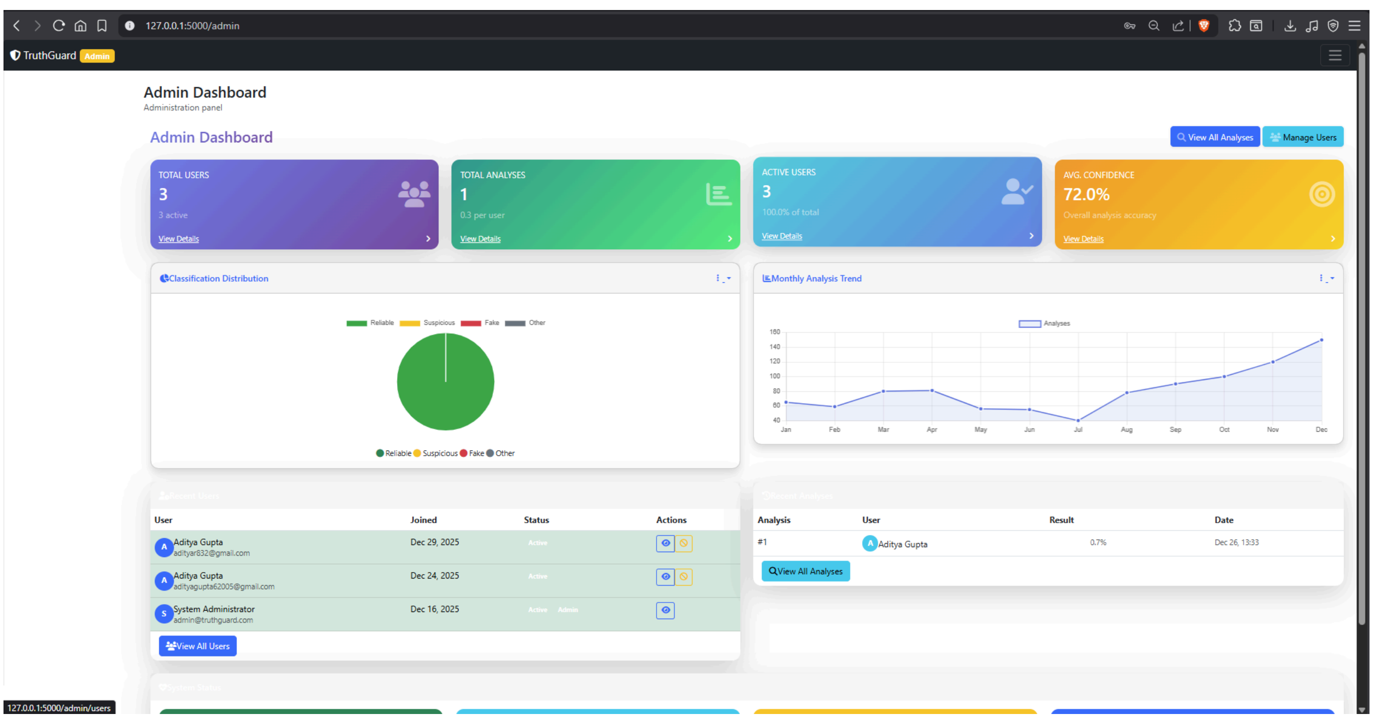

127.0.0.1:5000/admin
TruthGuard Admin
Admin Dashboard
Administration panel
Admin Dashboard
View All Analyses
Manage Users
TOTAL USERS
3
3 active
View Details
TOTAL ANALYSES
1
0.3 per user
View Details
ACTIVE USERS
3
100.0% of total
View Details
AVG. CONFIDENCE
72.0%
Overall analysis accuracy
View Details
Classification Distribution
Reliable Suspicious Fake Other
Monthly Analysis Trend
Analyses
Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
User Joined Status Actions
Aditya Gupta adityar832@gmail.com Dec 29, 2025
Aditya Gupta adityagupta62005@gmail.com Dec 24, 2025
System Administrator admin@truthguard.com Dec 16, 2025
View All Users
Analysis User Result Date
#1 Aditya Gupta 0.7% Dec 26, 13:33
View All Analyses
127.0.0.1:5000/admin/users

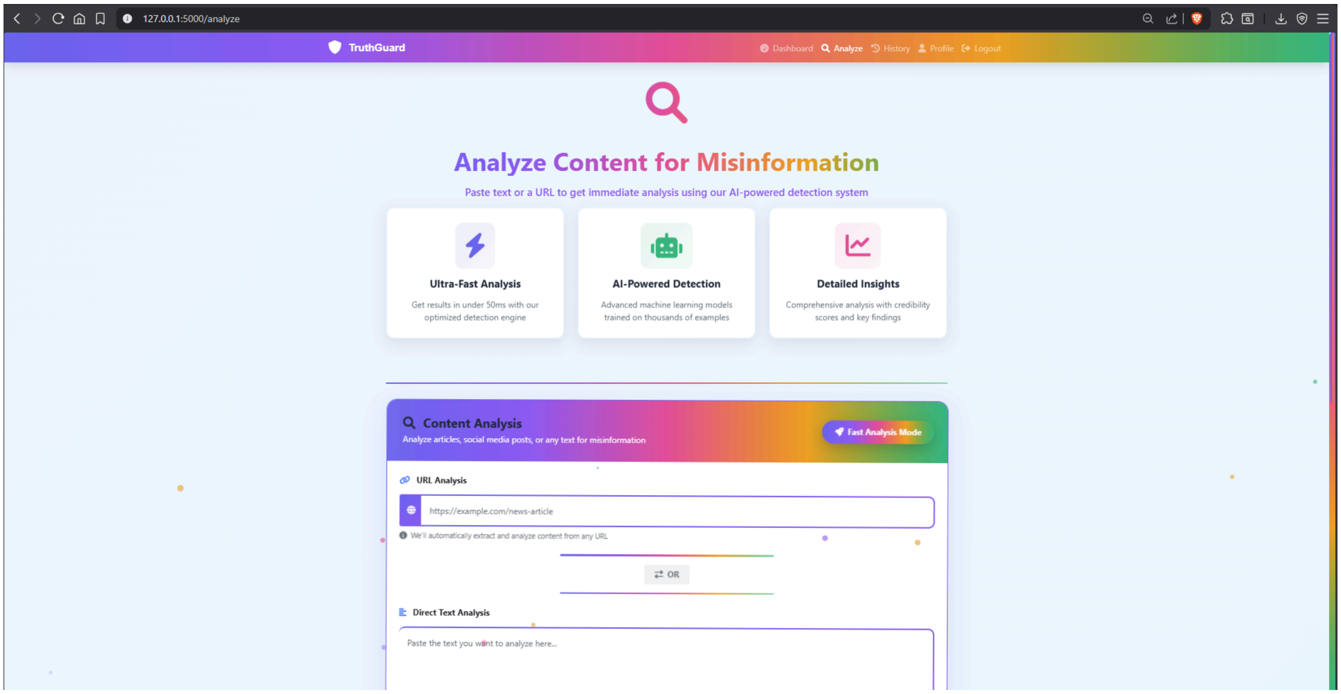
127.0.0.1:5000/analyze
TruthGuard
Dashboard Analyze History Profile Logout
Analyze Content for Misinformation
Paste text or a URL to get immediate analysis using our AI-powered detection system
Ultra-Fast Analysis
Get results in under 50ms with our optimized detection engine
AI-Powered Detection
Advanced machine learning models trained on thousands of examples
Detailed Insights
Comprehensive analysis with credibility scores and key findings
Content Analysis
Analyze articles, social media posts, or any text for misinformation
Fast Analysis Mode
URL Analysis
https://example.com/news-article
We'll automatically extract and analyze content from any URL
OR
Direct Text Analysis
Paste the text you want to analyze here...

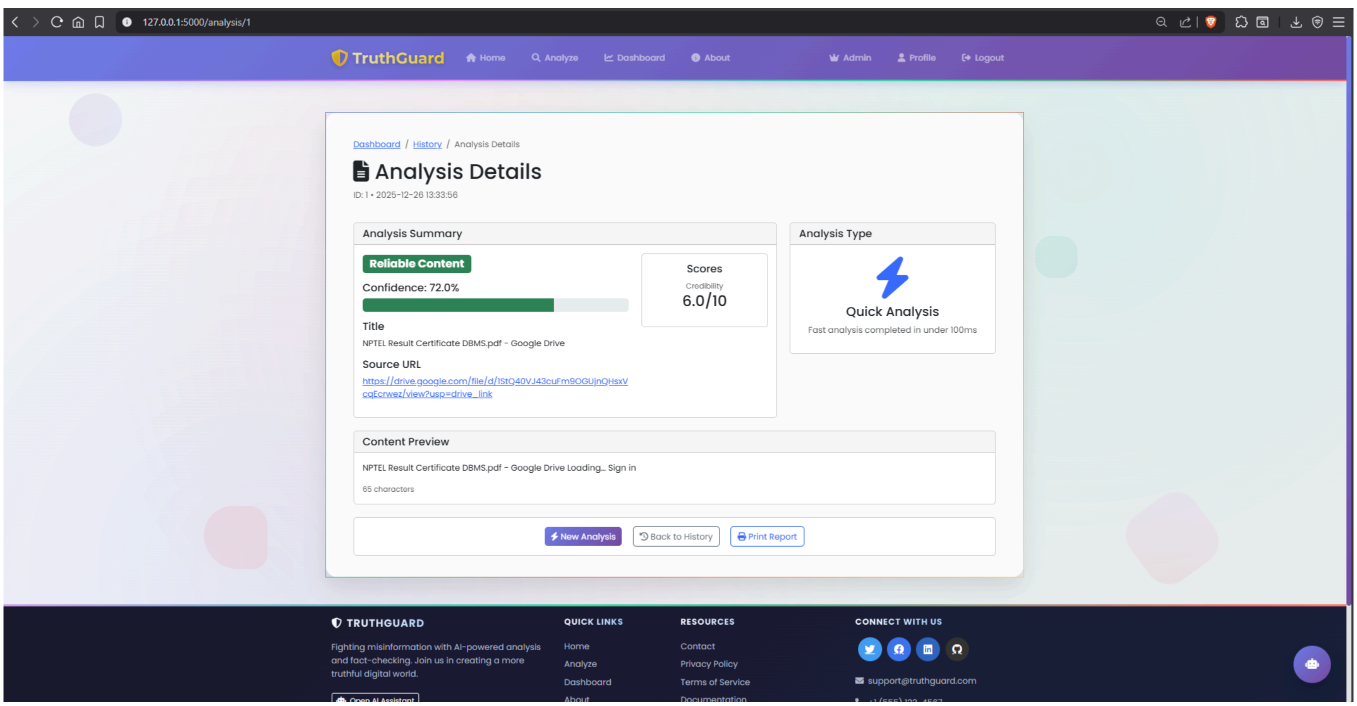
127.0.0.1:5000/analysis/1
TruthGuard
Home Analyze Dashboard About
Admin Profile Logout
Dashboard / History / Analysis Details
Analysis Details
ID: 1 • 2025-12-26 13:33:56
Analysis Summary
Reliable Content
Confidence: 72.0%
Scores
Credibility
6.0/10
Title
NPTEL Result Certificate DBMS.pdf - Google Drive
Source URL
https://drive.google.com/file/d/1StO40VJ43cuFm9OGUjnQHsxVcqEcrwez/view?usp=drive_link
Analysis Type
Quick Analysis
Fast analysis completed in under 100ms
Content Preview
NPTEL Result Certificate DBMS.pdf - Google Drive Loading... Sign in
65 characters
New Analysis
Back to History
Print Report
TRUTHGUARD
Fighting misinformation with AI-powered analysis and fact-checking. Join us in creating a more truthful digital world.
QUICK LINKS
Home
Analyze
Dashboard
RESOURCES
Contact
Privacy Policy
Terms of Service
CONNECT WITH US
support@truthguard.com

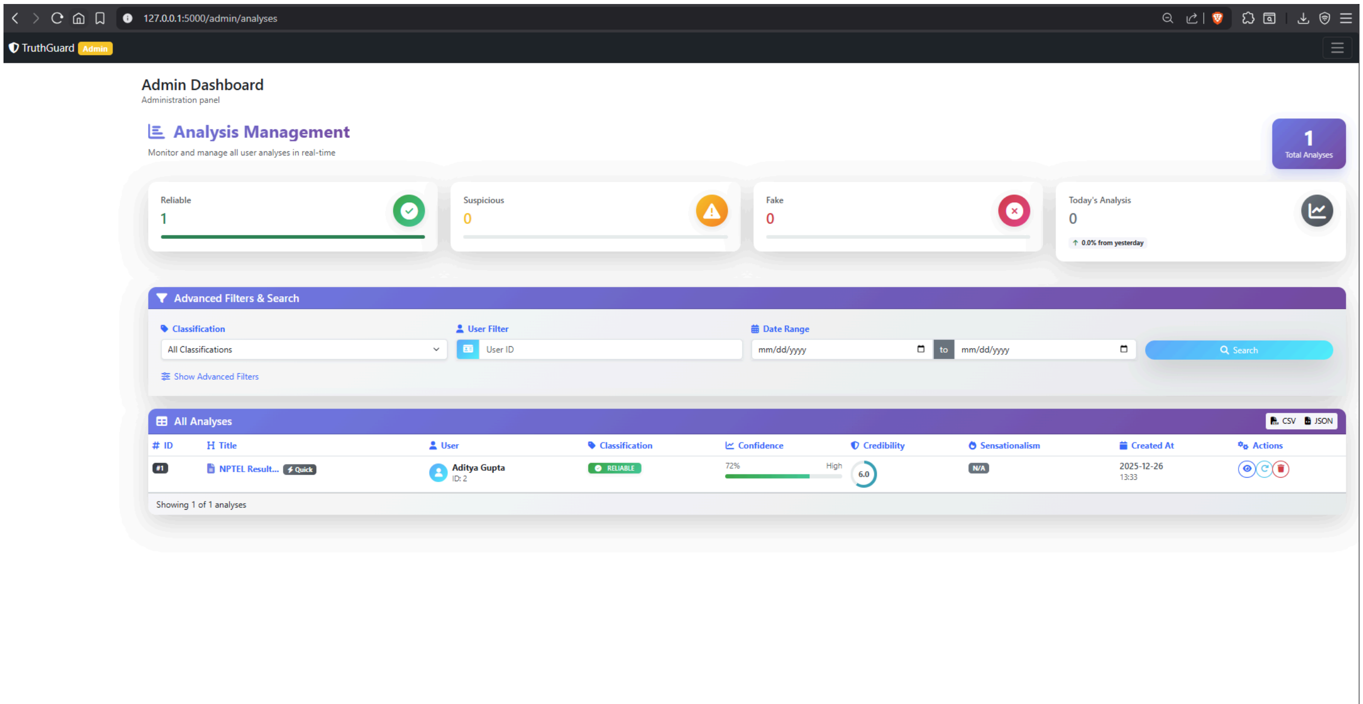
127.0.0.1:5000/admin/analyses
TruthGuard Admin
Admin Dashboard
Administration panel
Analysis Management
Monitor and manage all user analyses in real-time
1
Total Analyses
Reliable
1
Suspicious
0
Fake
0
Today's Analysis
0
0.0% from yesterday
Advanced Filters & Search
Classification
All Classifications
User Filter
User ID
Date Range
mm/dd/yyyy
to
mm/dd/yyyy
Search
Show Advanced Filters
All Analyses
CSV
JSON
ID
Title
User
Classification
Confidence
Credibility
Sensationalism
Created At
Actions
#1
NPTEL Result...
Quick
Aditya Gupta
ID: 2
RELIABLE
72%
High
6.0
N/A
2025-12-26
13:33
Showing 1 of 1 analyses

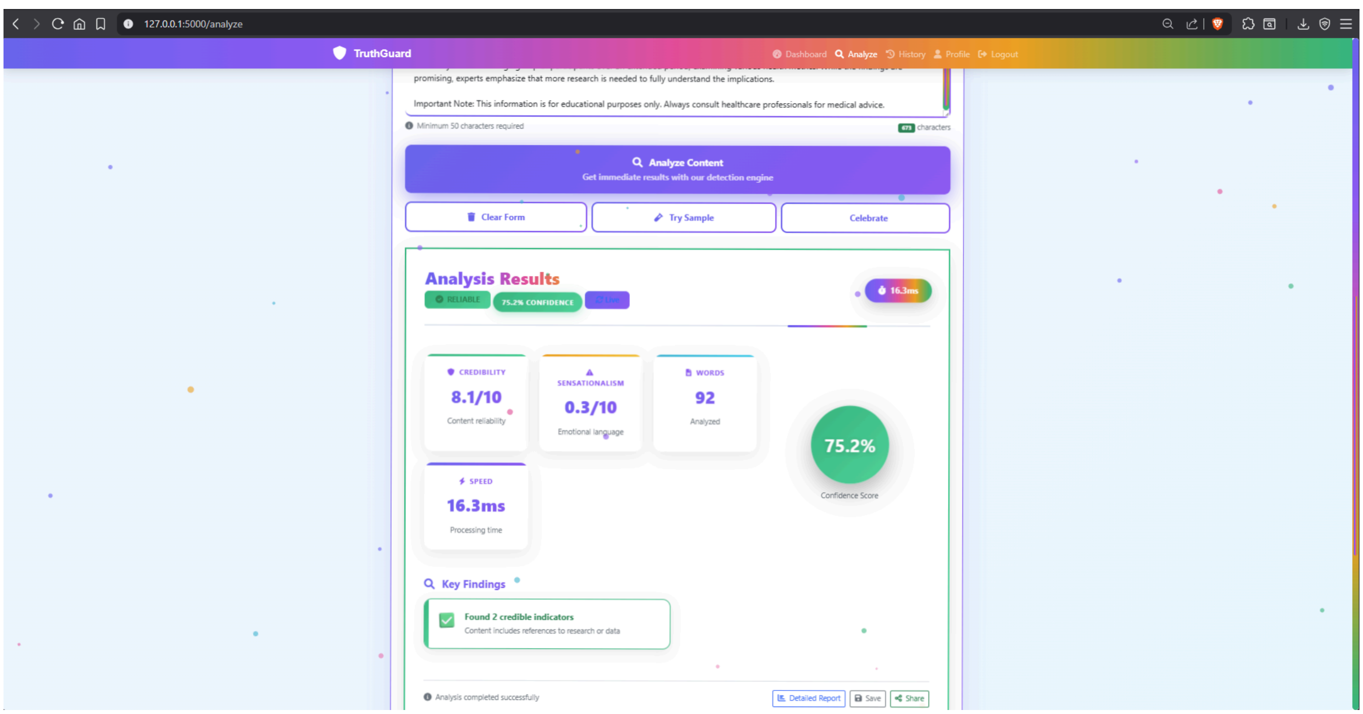
127.0.0.1:5000/analyze
TruthGuard
Dashboard
Analyze
History
Profile
Logout
promising, experts emphasize that more research is needed to fully understand the implications.
Important Note: This information is for educational purposes only. Always consult healthcare professionals for medical advice.
Minimum 50 characters required
characters
Analyze Content
Get immediate results with our detection engine
Clear Form
Try Sample
Celebrate
Analysis Results
RELIABLE
75.2% CONFIDENCE
16.3ms
CREDIBILITY
8.1/10
Content reliability
SENSATIONALISM
0.3/10
Emotional language
WORDS
92
Analyzed
SPEED
16.3ms
Processing time
75.2%
Confidence Score
Key Findings
Found 2 credible indicators
Content includes references to research or data
Analysis completed successfully
Detailed Report
Save
Share